

Homework 6

Alex Soupir

March 16, 2020

Packages: ISLR, class, MASS, mclust

Collaborators:

Answer all questions specified on the problem and include a discussion on how your results answered/addressed the question.

Submit your **.rmd** file with the knitted **PDF** (or knitted Word Document saved as a PDF). If you are having trouble with .rmd, let us know and we will help you, but both the .rmd and the PDF are required.

This file can be used as a skeleton document for your code/write up. Please follow the instructions found under Content for Formatting and Guidelines. No code should be in your PDF write-up unless stated otherwise.

There is no need for plotting in this assignment.

You do not need to include the above statements.

Please do the following problems from the text book ISLR or written otherwise.

1. Question 5.4.3 pg 198

We now review k -fold cross-validation.

- (a) Explain how k -fold cross-validation is implemented.

k -fold cross-validation is implemented by splitting the data into k sets, then removing one of the sets and training the model and testing on set. This is repeated for all k sets. Error is then calculated by averaging the MSE from all of the models.

- (b) What are the advantages and disadvantages of k -fold cross-validation relative to:

- i. The validation set approach?

Validation set is easy to implement, but can cause a bias based on which samples are included in the training and testing set, but the variation is much greater than that of the k -fold cross-validation.

- ii. LOOCV?

Leave one out can use a lot of computational power as mentioned in lecture. As shown on one of the slides, LOOCV may have slightly lower MSE than k -fold cross-validation but k -fold variation is low.

2. Question 5.4.5 pg 198 (use `set.seed(702)` to make results replicable)

In Chapter 4, we used logistic regression to predict the probability of `default` using `income` and `balance` on the `Default` data set. We will now estimate the test error of this logistic regression model using the validation set approach. Do not forget to set a random seed before beginning your analysis.

- (a) Fit a logistic regression model that uses `income` and `balance` to predict `default`.

- (b) Using the validation set approach, estimate the test error of this model. In order to do this, you must perform the following steps:
- Split the sample set into a training set and a validation set.
 - Fit a multiple logistic regression model using only the training observations.
 - Obtain a prediction of default status for each individual in the validation set by computing the posterior probability of default for that individual, and classifying the individual to the `default` category if the posterior probability is greater than 0.5.
 - Compute the validation set error, which is the fraction of the observations in the validation set that are misclassified.

```
## [1] 0.02791878
```

- (c) Repeat the process in (b) three times, using three different splits of the observations into a training set and validation set. Comment on the results obtained.

```
## [1] 0.02791878
```

```
## [1] 0.03002545
```

```
## [1] 0.02304147
```

The error is very consistent between the different splits of the data. This is the error of those that are misclassified

- (d) Now consider a logistic regression model that predicts the probability of `default` using `income`, `balance`, and a dummy variable for `student`. Estimate the test error for this model using the validation set approach. Comment on whether or not including a dummy variable for `student` leads to a reduction in the test error rate.

```
## [1] 0.02791878
```

Adding the variable for `student` produces similar (same) error when setting the seed to the same one as used in (b). This suggests that the dummy variable doesn't reduce the error test rate when included.

3. Question 5.4.7 pg 200

In Sections 5.3.2 and 5.3.3, we saw that the `cv.glm()` functions can be used in order to compute the LOOCV test error estimate. Alternatively, one could compute those quantities using just the `glm()` and `predict.glm()` functions, and a for loop. You will now take this approach in order to compute the LOOCV error for a simple logistic regression model on the `Weekly` data set. Recall that in the context of classification problems, the LOOCV error is given in (5.4).

- Fit a logistic regression model that predicts `Direction` using `Lag1` and `Lag2`.
- Fit a logistic regression model that predicts `Direction` using `Lag1` and `Lag2` *using all but the first observation*.
- Use the model from (b) to predict the direction of the first observation. You can do this by predicting that the first observation will go up if $P(\text{Direction} = \text{"Up"} | \text{Lag1}, \text{Lag2}) > 0.5$. Was this observation correctly classified?

```
##      1
```

```
## "Up"
```

```
## Actual Direction:
```

```
## [1] Down
```

```
## Levels: Down Up
```

The observation was incorrectly classified.

- (d) Write a for loop from $i = 1$ to $i = n$, where n is the number of observations in the data set, that performs each of the following steps:
 - i. Fit a logistic regression model using all but the i th observation to predict `Direction` using `Lag1` and `Lag2`.
 - ii. Compute the posterior probability of the market moving up for the i th observation.
 - iii. Use the posterior probability for the i th observation in order to predict whether or not the market moves up.
 - iv. Determine whether or not an error was made in predicting the direction for the i th observation. If an error was made, then indicate this as a 1, and otherwise indicate it as a 0.
- (e) Take the average of the n numbers obtained in (d)iv in order to obtain the LOOCV estimate for the test error. Comment on the results.

```
## [1] 0.4499541
```

The error is rather high. Looking at the values of individual errors and comparing them with the `Direction` variable in the data, the model doesn't seem to be predicting the same value for all observations, meaning it isn't predicting "Up" or "Down" for each observation. Interestingly though, if it were to predict Up for all of the observations, the error (the observations that are incorrect over the total number of observations) would be $484/1089$, which is 0.444 and is lower than that of the model's prediction error.

4. Write your own code (similar to Q #3. above) to estimate test error using 6-fold cross validation for fitting linear regression with $mpg \sim horsepower + horsepower^2$ from the Auto data in the ISLR library. You should show the code in your final PDF.

Going off of the slide for 5.1.3 k-fold Cross Validation from lecture, I will split the data into k (6) different sets. Then, I'll remove one set, train the model on the remaining data, and test on the set that was removed, calculating the MSE. I'll do this for all k data sets, storing the MSE for each as a term in a vector. After all k sets are tested, I will take the average of all of the MSE terms.

```
set.seed(702)
data("Auto")
error = double()
sample.ind = sample(6, nrow(Auto), replace=true)
for(i in seq(max(sample.ind))){
  train.4 = Auto[c(sample.ind != i),]
  valid.4 = Auto[c(sample.ind == i),]

  model.4 = glm(mpg ~ horsepower + horsepower^2, data=train.4)
  preds.4 = round(predict(model.4, valid.4), digits=0)
  error[i] = mean((preds.4 - valid.4$mpg)^2)
}
mean(error)
```

```
## [1] 24.67371
```

5. Last homework you started analyzing the dataset you chose from [this website](#). Now continue the analysis and perform Logistic Regression, KNN, LDA, QDA, MclustDA, MclustDA with EDDA if appropriate. If it is not possible to perform any of the methods mentioned above please justify why.

The data has 6 columns: BIRADS is the assessment provided by double-review process with physicians with 1 being benign and 5 being most likely malignant, Age of the patient, Shape of the mass which is a factor and not integer, Margin of the mass which is also nominal, Density of the mass which is most likely a factor based on the description [high=1, iso=2, low=3 fat-containing=4], and Severity of the mass with 0 being benign and 1 being malignant.

Since there are a fair amount of observations with na, I'll remove those that are not complete cases. Even though BIRADS is not a predictor, I want to make sure that the values are there for comparison in the future.

Logistic Regression - LOOCV

```
## [1] 0.186747
```

KNN

```
## Error rate for KNN (seed 702) (K =1): [1] 0.2590361
## Error rate for KNN (seed 702) (K =2): [1] 0.2409639
## Error rate for KNN (seed 702) (K =3): [1] 0.2289157
## Error rate for KNN (seed 702) (K =4): [1] 0.2289157
## Error rate for KNN (seed 702) (K =5): [1] 0.2228916
## Error rate for KNN (seed 702) (K =6): [1] 0.2409639
## Error rate for KNN (seed 702) (K =7): [1] 0.2349398
## Error rate for KNN (seed 702) (K =8): [1] 0.2409639
## Error rate for KNN (seed 702) (K =9): [1] 0.246988
## Error rate for KNN (seed 702) (K =10): [1] 0.2409639
## Error rate for KNN (seed 702) (K =11): [1] 0.2349398
## Error rate for KNN (seed 702) (K =12): [1] 0.2349398
## Error rate for KNN (seed 702) (K =13): [1] 0.2349398
## Error rate for KNN (seed 702) (K =14): [1] 0.2349398
## Error rate for KNN (seed 702) (K =15): [1] 0.2349398
## Error rate for KNN (seed 702) (K =16): [1] 0.2349398
## Error rate for KNN (seed 702) (K =17): [1] 0.2349398
## Error rate for KNN (seed 702) (K =18): [1] 0.2349398
## Error rate for KNN (seed 702) (K =19): [1] 0.2349398
## Error rate for KNN (seed 702) (K =20): [1] 0.2349398

## Minimum error was produced using k=5 with an error of 0.2228916
```

LDA

```
## Error rate for LDA on test split (seed 702):
## [1] 0.2349398
```

QDA

```
## Error rate for QDA on test split (seed 702):
## [1] 0.2349398
```

MclustDA

```
##      Predicted
## Class    0    1
##      0 215 126
##      1  23 300

## Error rate for MclustDA on test split (seed 702):
## [1] 0.2243976
```

MclustDA with EDDA

```
##      Predicted
## Class    0    1
##      0 253  88
##      1  43 280

## Error rate for MclustDA on test split (seed 702):
## [1] 0.1972892
```

Doctors comparison

To compare to what the double-review by doctors predicted, I decided to look at their assessment. To do this, I first looked at the summary of the 2 factor variables BIRADS and Severity. This shows that there are more actual benign diagnoses than malignant, however doctors weighed more towards a diagnosis of malignant. This may be because it's better to be safe than sorry in diagnoses and determining a mammogram mass as benign when it is malignant can be dangerous. Next, I created a confusion matrix where I converted BIRADS assessment values of 4 and 5 to malignant and 1 through 3 as benign. The error of these predictions was incredibly high due to the doctors pushing towards malignant. However, only converting an assessment of 5 to malignant lowered the error from 0.49 (BIRADS 4 and 5 as malignant) to 0.17 (BIRADS 5 as malignant).

These errors are on both extremes of the methods tested above, with the BIRADS 4 and 5 having at least 2x the error than any of the methods and BIRADS 5 having lower error than any of the methods. Of the methods tested above, the logistic regression produced the lowest error with LOOCV (0.186747). Following this, the next lowest error was produced with KNN where $k=5$ (0.2228916), then the MclustDA with EDDA model type (0.1972892) and MclustDA default model type was slightly higher (0.2243976). QDA (0.2349398) and LDA (0.2349398) were the worst in terms of error in their predictions.

```
## Summary of the BIRADS factor by doctors:
##    0    2    3    4    5
##    5    7   24 468 316

## Summary of the actual diagnoses:
##    0    1
## 425 395

##
## Confusion matrix of the BIRADS assessment by doctors (>3 is malignant):
```

```

##      docs
##      0   1
##    0 29 396
##    1   7 388

## Error of doctor prediction (BIRAD>3 = Malignant):
## [1] 0.4914634

##
## Confusion matrix of the BIRADS assessment by doctors (>4 is malignant):
##      docs
##      0   1
##    0 394  31
##    1 110 285

## Error of doctor prediction (BIRAD>4 = Malignant):
## [1] 0.1719512

```

References

- ISLR Book
- Stackoverflow
-